

Which Hand-Motion Composition Axes Are Worth Measuring

Anonymous authors (double-anonymous review)

Abstract—Modular motion priors decompose hand motion into a bank of primitives, on the premise that novel recombinations remain usable, which is rarely measured. We build an instrument for measuring it: a paired split scoring a naive against an informed arm on held-out compositions, five gates that reject sweeps whose numbers cannot mean what they claim, and a control whose true penalty is zero. Read against that control, the penalty follows one rule over eight axes of three public datasets: every axis carrying compositional difficulty has positive interaction between its two factors, and no axis without it does. Three axes could have falsified the second half, and all three sit at the control ($p = 0.60$ to 0.96). A statistic on raw pose arrays finds the interaction in minutes, so a non-positive reading is a cheap reason to skip an axis; it licenses nothing in the other direction and ranks axes badly, its association with the penalty falling from $r = 0.843$ to 0.463 . Where difficulty exists a modular bank pays a smaller penalty than a parameter-matched monolithic latent, but at matched information rate the control reproduces 90% of the normalized advantage, so we claim nothing for modular structure. On a Shadow Hand, sixteen of seventeen defined comparisons fail to separate novel from familiar compositions.

Index Terms: Dexterous manipulation, learning from demonstration, motion primitives, deep learning methods, evaluation and benchmarking.

I. INTRODUCTION

Motion priors for dexterous hands are usually monolithic, emitting one continuous latent per control step. LAMP [1] is the current instance for hands, and it rejects a quantized alternative because “codebook switches can turn small latent changes into abrupt hand-command jumps during contact-rich residual RL”. Whole-body animation went the other way. MotionBricks [2] builds a modular latent backbone over more than 350,000 clips, but a human composes its primitives through an authoring interface, so it does not test whether the units of a *learned* bank mean anything stable.

Both architectures assume that separately learned primitives recombine and that novel recombinations stay usable. Whether modularity buys that at the scale hand datasets offer is not obvious. MotionBricks trains on 10^5 clips; the largest public hand dataset used here provides 10^3 . Whether any of it reaches a robot is a further question, as a composition can be valid human motion and still unreachable on a hand whose fingers are spaced differently. We build a measurement for the first and report a null for the second; our contributions are as follows.

- 1) A rule for which composition axes are worth measuring: compositional difficulty appears only where the two factors interact, over eight axes on three datasets, with three axes able to falsify it and none doing so (Section V-B). A statistic computed on raw poses in minutes finds the interaction. It is a reason to skip an axis, not to sweep one, and it ranks axes badly.

- 1) An instrument for measuring compositional difficulty: a paired split, five gates that reject sweeps whose numbers cannot mean what they claim (Section III-C, Table I), and a control whose true penalty is zero, against which every penalty is read. Four gates exist because an earlier design returned plausible numbers it could not have earned, and the control matters because a test against zero at forty seeds detects the instrument’s own bias.
- 1) A null on the hardware question, and an account of why it is a weak null. Three things weaken it: the retargeting it rests on has no collision term, the banks behind it were trained on the one dataset our own instrument finds null, and the equivalence test that would license “the same” fails its own sensitivity check.

That transfer is our measuring apparatus rather than our contribution. Work carrying human hand-object motion onto real hardware [7], [13] is further along that axis than this paper, and the three modular architectures we compare are likewise prior art.

II. RELATED WORK

LAMP [1] learns a history-conditioned continuous latent prior over hand actions and uses it as the action space for real-world residual reinforcement learning. DexSynRefine [7] represents human hand-object motion with motion-manifold primitives and grounds it in physics with task-space residual RL. Neither evaluates whether a prior’s primitives generalize to combinations absent from training. Our parameter-matched baseline *is* a LAMP-style continuous latent, and our modular arm takes its backbone from MotionBricks [2], whose primitives a human authors through an interface; what we study is the learned bank that replaces them.

Modular motion models recur in mechanically distinct forms, from an authored-primitive latent backbone [2] to switching linear [3] and nonlinear [4] dynamical systems and mode-adaptive gating [5], all sharing the premise that separately learned primitives can be recombined. A survey of over 700 movement-primitive papers [8] lists composing primitives and evaluating frameworks among its open questions.

Recent work uses “compositional” in senses that do not answer that question: CoDex [9] composes task stages and learns no motion prior, UniCross [10] composes hand-defined skills, and object-centric neural fields [11] mix object-conditioned primitives, all holding out objects rather than the pairing of two factors. Adversarial Concept Search [12] predicts failing concept combinations but needs a trained




Fig. 1. The same 90%-flexion reference pose projected onto the anatomical model from each dataset, after the conventions are inferred per joint; thumb red, index orange, middle yellow, ring green, little blue. Before the fix an assumed abduction axis put 100% of GRAB and DexYCB frames outside thumb_mcp_abd’s range.

model; SPIDER [13] retargets demonstrations at scale, and we claim no contribution there.

III. METHOD

A. Hand representation

Motion is 27 anatomical degrees of freedom in native units: four per finger (MCP flexion and abduction, PIP, DIP), five for the thumb, six for the wrist. Normalization into $[-1, 1]$ comes from the joint-limit box, never from dataset statistics, which keeps a prior trained on one source meaningful on another.

Public hand datasets provide MANO [20] parameters, and projecting them onto an anatomical model requires three conventions: which Euler axis carries flexion, its sign, and which of the remaining two carries abduction. We infer all three per joint and assume none. We did assume them first, and the thumb and the fingers turned out not to follow each other’s convention.¹ Assuming costs in both directions, and we paid both. Take abduction as $(\text{flexion} + 1) \bmod 3$, correct for the fingers, and 100% of GRAB and DexYCB frames fall outside thumb_mcp_abd’s 20° range. Pool PIP variance across fingers instead, and their differing resting offsets count as flexion; on OakInk2 that selected the wrong axis and pinned index_pip_flex at its limit in every frame. Inferring per joint halves that dataset’s projection residual and takes its PIP joints from 37.3% inside their anatomical range to 99.9%. Fig. ?? shows the three datasets agreeing on a reference pose once the conventions are inferred.

B. The two priors

Architecture. Both arms encode a 32-frame window of 27 anatomical DOF and reconstruct it under a bounded output, which makes joint-limit satisfaction architectural. The modular arm emits a relaxed categorical assignment over a bank of K primitives, one per frame, each a 32-dimensional embedding, plus one 4-D style code for the window. That code is kept small so it cannot let the decoder bypass the bank. The assignment uses Gumbel-softmax [22] in training and a plain softmax at evaluation, so the reported numbers pass through a convex combination and never a hard switch.

¹MANO’s zero pose is a flat hand chosen for compatibility with SMPL [20], and its release carries a kinematic tree with no per-joint rest rotations, so every joint frame is parallel to the template’s. The thumb’s is therefore no better aligned to its own anatomy than a finger’s, and its flexion and abduction fall on different Euler components.

The baseline replaces the bank with a 12-D Gaussian latent, also per frame. Both encode with bidirectional two-layer GRUs of 256 units and decode with unidirectional ones, and their parameter counts agree to within 0.6% at the $K = 12$ that Section V-A’s sweeps use; the release gives the counts.

Objective. Both are trained as beta-VAEs [23] on windowed reconstruction, with mean-squared error on pose, a first-difference term because the data is 30 fps motion, and β times a KL with a free-bits floor of 0.02 nats per dimension. The weight β is annealed over the first epochs, and validation is always scored at its final value. The modular arm adds four terms that shape the bank rather than the reconstruction: a switch penalty between adjacent frames, a per-frame entropy term, a usage-entropy reward and a load-balance term. A fifth is optional and off in Section V-A’s sweeps, and it penalizes the spread of the velocity profiles a primitive produces when decoded after several predecessors and from several initial poses, velocity rather than pose because pose is dominated by the initial condition. It optimizes the quantity Section V-C measures, which is why Section V-C reads its own number as a target being met.

Rate matching. The quantity to match is the rate of information flow, which the parameter count does not fix. A window-level latent would push 32 frames through one code while the bank emits an assignment every frame, so the baseline gets a per-frame latent too. On nominal capacity that leaves the baseline ahead, as a 12-D continuous Gaussian carries more per frame than a 12-way categorical whose 4-D style code is paid once per window. On realized rate, by contrast, the arms are not matched at all, and the gap favors the modular arm. The confirmatory axis gives 0.207 nats per frame for the baseline against 2.275 for the assignment channel, and two further sweeps give roughly elevenfold as well. The release shows that the baseline’s KL penalty is nonetheless binding, that its latent is not collapsed, and that the assignment is nearly deterministic given the input, which is what makes it a high-rate channel.

The matched arm. Raising the baseline’s free-bits floor to 0.186 per dimension closes the gap. Its realized KL rises to 1.846 nats per frame against the modular arm’s 2.312, a ratio of 1.25, where the unmatched baseline in the same sweep sits at 0.206 for a ratio of 11.2. The confirmatory axis was re-swept on that basis, with all three arms trained together over the same 70 declared seeds, all of which returned every arm. That intervention produces a result, so Section V-A reports it, with the sensitivity check belonging to the intervention itself, given with the release. One caution applies to any comparison drawn across sweeps, and Section VI gives it.

C. Paired compositional measurement

We score the same target trajectories under two priors that differ only in whether their training data contained the target’s composition,

$$\begin{aligned} \text{penalty} &= \text{MSE_naive}(T) - \text{MSE_informed}(T) \\ \text{paired} &= \text{penalty_baseline} - \text{penalty_modular} \end{aligned}$$

Both arms see the same number of training trajectories and are scored on identical test data, so trajectory difficulty

Gate	What it caught
primitive collapse	primitives_used = 1.0, a modular model that is monolithic in disguise
informed coverage	held compositions absent from the informed set, which pins the penalty at zero by construction
composition leak	targets whose exact fine composition sat in the informed training set, against none for naive
DOF health	thumb_mcp_abd pinned at its limit in every GRAB and DexYCB frame under an assumed abduction axis
split balance	guard only; the risk that an unpaired design measures set difficulty rather than compositional difficulty

TABLE I
THE FIVE GATES, AND WHAT EACH CAUGHT

and set size difference out and the arms’ evaluation noise is correlated. On the corrected bundles the imbalance is -6.1% to +10.7%, mean +1.4%, and the gate passes on that mean. It passes on a statistic that is not the concern, though. The corrected per-seed spread is wider than the pre-fix one, so the risk the unpaired design runs is present in the release, and pairing removes it by construction.

The fifth gate guards a risk the corrected data does not show. The other four each caught a defect in an earlier design, and two still exit 1 on the released artifacts, informed coverage on the confirmatory and affordance axes and DOF health at 40% and 33% against its own 25% threshold. The release gives the transcript each gate was read from, the sweeps each was run over, and the breadth rule Section V-B applies. The DOF gate tests per-DOF pinned fraction and variance directly. It replaces a clip-fraction statistic, and the reason matters. That statistic is defined on values *before* clamping, so an adapter that passes it post-clamp values returns ≈ 0 whatever the data does.

D. Predicting which axes are worth measuring

A full sweep costs hours of GPU per axis, so a cheap statistic decides beforehand which axes repay that cost, and most axes do not. The compositional penalty asks whether knowing the *pairing* helps beyond knowing each factor separately, which is an *interaction*, so we fit two nested least-squares models per axis on PCA-reduced windowed poses,

additive pose $\sim$ left + right

full pose $\sim$ left + right + cell

interaction $R^2(\text{full}) - R^2(\text{additive})$

with a permutation null that shuffles labels over whole trajectories, since 30 fps windows within one trajectory are near-duplicates. A variance-explained statistic on the composition label alone does not work. Of the three real datasets scored on η^2 , GRAB’s excess (+0.0538) and OakInk-Image’s (+0.0547) are indistinguishable, yet GRAB’s swept axis returns no compositional difficulty ($p = 0.494$) where OakInk-Image’s returns +0.00982 on the id-prefix axis Section IV admits only for comparison. The interaction term separates the two, at $z = 8.3$ on the confirmatory axis against 3.9 on GRAB’s. Structure alone is not the reason either. The number of transitions a leak-free split can hold out, which the split builder computes for free, anticorrelates with the measured penalty over the swept axes ($r = -0.585$) and ranks OakInk2’s

transition axis first, where the sweep returns nothing; the interaction statistic ranks it last.

E. Retargeting to the Shadow Hand

Robot model. The robot is the Shadow Hand E3M5 model from MuJoCo Menagerie [18], in MuJoCo [17]. Anatomical DOF map to its joints by role, corroborated by an independent check. The two models’ ranges were designed separately, so their agreement is evidence that the roles line up.² We transfer angles directly and clip them; rescaling between the limit boxes would hide the saturation this is meant to measure.³

Contact and penetration. Fingers resting together in a closed grasp register contacts and are valid poses, so the report separates *touching* from *penetrating beyond a tolerance* and counts only the second. At zero configuration the model has zero contacts and every geometry margin is 0, so measured penetration reflects real geometry, although that check is not in the release.

Abduction correction. We run coordinate descent on the abduction joints only, as flexion carries what the motion *is* and abduction carries how far apart the fingers are held, which is what the index-against-middle failure needs more of. Each step tries every abduction joint at $\pm 1.5^\circ$ and keeps the move that most reduces the deepest penetration, and which sign spreads which finger is read from the model.

IV. EXPERIMENTAL SETUP

We use four public datasets (OakInk-Image [14], OakInk2 [15], DexYCB [16], GRAB [6]) plus two synthetic controls, one with planted via-point compositional difficulty and one purely interpolable. GRAB uses 8 of its 10 subject archives, which is how many downloaded. Sweeps run at a budget of 256 training trajectories per arm for 120 epochs, scored by a paired per-seed *t*-test with a sign test alongside, since a *t*-test alone cannot rule out a few extreme seeds carrying an effect. A result is declared only when both clear $\alpha = 0.05$, which is what “clears both declared tests” and “reaches neither” refer to below. The banks of Sections V-C and V-D use 200 epochs at stride 8 on full data.

The interpolable control returns no compositional advantage: +0.00154 with 12 of 20 ($p = 0.154$) against the unmatched baseline and +0.00155 with 13 of 20 ($p = 0.090$) against the rate-matched one, at budget 256 over twenty seeds. Its difficulty, +0.00258 ($p = 0.0002$), is the floor Sections V-B and VI read every penalty against.

Composition axes come from each dataset’s own taxonomy where one exists; OakInk ships object categories, functional classes and affordance attributes, and all 100 object ids resolve. An earlier axis used the first character of the object id, which records a sub-collection, and is carried through for comparison alone.

²index_mcp_abd [-20, 20] against FFJ4 [-20, 20], thumb_mcp_abd [-10, 10] against THJ3 [-12, 12], thumb_ip_flex [-15, 80] against THJ1 [-15, 90]. Where they agree less well (thumb_mcp_flex [-10, 55] against THJ2 [-40, 40]) the per-joint clip rate reports it.

³The Menagerie model has 24 joints but 20 actuators, each finger’s two distal joints sharing a tendon, so writing qpos commands PIP and DIP independently where the physical hand cannot.

data / axis	advantage (wins)	d	p	difficulty
synthetic, via-point	+0.00491 (17/18)	1.13	0.0002	+0.00904
OakInk-Image, category $\times$ intent	+0.00397 (48/69)	0.41	0.0012	+0.01042
OakInk-Image, class $\times$ intent	+0.00873 (9/12)	0.65	0.046	+0.01650
OakInk-Image, affordance $\times$ intent (exploratory)	+0.00365 (86/130)	0.33	0.0002	+0.00873
OakInk-Image, id-prefix $\times$ intent (deprecated)	+0.01022 (60/70)	0.91	<0.0001	+0.00982
GRAB, shape $\times$ fine intent	+0.00305 (25/40)	0.18	0.269	+0.00041
OakInk2, annotated transitions	-0.00010 (4/12)	-0.05	0.860	+0.00082
GRAB, shape $\times$ intent class (post-submission)	+0.00185 (25/40)	0.25	0.129	+0.00103
OakInk2, scene $\times$ verb (post-submission)	-0.00097 (18/40)	-0.21	0.186	+0.00051
synthetic, interpolated (control, true zero)	+0.00154 (12/20)	0.33	0.154	+0.00258

TABLE II
THE MODULAR ADVANTAGE, AND THE DIFFICULTY IT ACTS ON

V. RESULTS

A. The modular arm’s advantage, and what rate matching does to it

The advantage is how much smaller the modular arm’s compositional penalty is than the baseline’s, and how often. As Table II shows, it is positive on six of the seven interpretable axes, the six real ones and the synthetic control. It clears Section IV’s declared pair of tests on four; a fifth, the class axis, clears the t -test alone. The two exceptions are the axes whose difficulty is itself indistinguishable from zero: GRAB’s shape axis, positive but at $p = 0.269$, and OakInk2’s, the one negative value in the table. The difficulty the advantage acts on is a separate quantity, and we report the advantage first as the smaller and less certain of the two. The last column reads as the size of the problem and the second as how much of it the modular arm removes, so on the confirmatory axis +0.00397 of +0.01042 is 38%, and DexYCB’s sweep is omitted for the reason Section V-B gives.

On that confirmatory axis the modular arm wins 48 splits in 69, seven in ten, which is the number to weigh before the effect size. It recurs on the class axis at $d = 0.65$, swept after the screen mis-ranked it (Section V-B), where the sign test is 9 of 12 at $p = 0.146$ and falls short of Section IV’s bar. We report that sign test because Section VI treats the same statistic as disqualifying on the affordance axis.

Scale is the first thing to rule out. The baseline’s absolute error is about a third higher than the modular arm’s, so a larger absolute penalty could be an artifact of error scale. Normalizing each seed’s penalty by its own naive error removes that confound, and the advantage then survives on two of the three official axes: +0.076 with 8 of 12 on the class axis, clearing neither test ($t p = 0.118$); +0.045 with 77 of 130 on the affordance axis, clearing both ($t p = 0.004$, sign $p = 0.043$); and +0.046 with 39 of 69 on the confirmatory category axis, clearing the t -test ($p = 0.021$) and not the sign test ($p = 0.336$), so normalizing costs that axis its

win rate, from 48 of 69 raw. The normalized metric is a robustness check that failed: two axes with no measurable difficulty return more on it than the confirmatory axis, GRAB’s shape axis at +0.116 and the withdrawn DexYCB axis at +0.043 against +0.046. Wherever the modular arm reconstructs worse, dividing by naive error manufactures an advantage; the release gives that audit. The category axis carries the confirmatory claim and the affordance axis is exploratory, for reasons Section VI gives.

Rate matching. That normalization also decides how Section III-B’s rate-matched comparison reads, and the two metrics read differently. On the scale-normalized penalty, whose premise is met here because the rate-matched baseline reconstructs better than either other arm, matching appears to remove nothing: the advantage is +0.104 against the unmatched baseline and +0.111 against the matched one, 52 of 70 at $d = 0.54$, clearing both declared tests. That reading does not survive the negative control. On *pc_easy*, whose paired advantage is zero by construction, the same normalized contrast returns +0.099 at 18 of 20 and clears both tests, 90% of the figure just quoted. Normalization divides by each arm’s naive error, and on the control the matched baseline reconstructs 4.04 times better than the modular arm against 1.41 on real data, so dividing inverts the error-scale confound instead of removing it. On the raw penalty the same matching removes most of the effect, from +0.00693 at 50 of 70 wins and $d = 0.61$ to +0.00254 at 41 of 70 and $d = 0.23$, reaching neither declared test, with a raw floor of +0.00155 on the same control and arm, 61% of it. One metric loses the effect and the other produces almost the same effect on data that has none, so we claim no modular advantage at matched rate.

Other decompositions. Two off-the-shelf decompositions, learned velocity fields under the same per-frame assignment (the switching formulation of [3], [4]) and MANN’s per-frame expert blending [5], also beat the parameter-matched baseline on synthetic data, +0.00567 and +0.00968 over 8 seeds ($p = 0.001$ and 0.010) against the bank’s +0.00491 over 18 ($t p = 0.0002$), and carry the same rate confound.

B. Which axes carry compositional difficulty

The screened axes. We screened fifteen axes, of which the eight in Table III are those whose sweeps are interpretable; the other seven are unswept, four of them because no leak-free split exists (below), and are listed with the release. One, DexYCB’s subject $\times$ shape, was also swept and is withdrawn below. Two of the eight, GRAB’s shape $\times$ intent class and OakInk2’s scene $\times$ verb, were swept after the initial submission under a protocol pre-registered with the release. Left and right give the two factors’ separate η^2 , excess is the interaction term of Section III-D minus its permutation-null mean with the z beside it, and the last column repeats Table II’s difficulty.

Informative factors are not enough on their own. OakInk2’s two factors are the most informative in Table III, at 0.188 and 0.243 against OakInk-Image’s 0.097 and 0.045, left factor first in each pair, and pairing them adds nothing

dataset, axis	left	right	excess (z)	difficulty
OakInk-Image, affordance $\times$ intent	0.054	0.045	+0.0309 (14.9)	+0.00873
OakInk-Image, category $\times$ intent	0.097	0.045	+0.0238 (8.3)	+0.01042
OakInk-Image, id-prefix $\times$ intent (deprecated)	0.028	0.045	+0.0182 (12.9)	+0.00982
GRAB, shape $\times$ fine intent	0.028	0.044	+0.0063 (3.9)	+0.00041
OakInk-Image, class $\times$ intent	0.022	0.045	+0.0061 (4.1)	+0.01650
OakInk2, annotated transitions	0.188	0.243	-0.0120 (-2.6)	+0.00082
GRAB, shape $\times$ intent class	0.028	0.009	-0.0005 (-0.4)	+0.00103
OakInk2, scene $\times$ verb	0.056	0.175	-0.0001 (-0.1)	+0.00051

TABLE III
THE INTERACTION SCREEN AGAINST THE MEASURED PENALTY

measurable, while GRAB’s swept axis is less informative on both and also null. The axes carrying difficulty share the interaction term and the informative ones lack it. A practitioner choosing what to record should therefore measure the interaction. It is the harder property to arrange, and the one a factor-by-factor look at a dataset will miss.

We read “carries difficulty” against the instrument’s own control rather than against zero. The interpolable control returns +2.9% of the naive error over twenty seeds (sd 2.8), significant against zero at $p = 0.0002$ (Section VI): at the sample sizes used here a test against zero detects the paired design’s own bias. Against that control, every axis with positive interaction that carries difficulty lies ten standard errors above it (category +15.2%, affordance +12.2%, Welch $p < 10^{-12}$), and every axis without positive interaction is indistinguishable from it. Three such axes could have shown otherwise. OakInk2’s annotated transitions (excess -0.0120) return +3.4% over twelve seeds ($p = 0.79$ against the control); GRAB’s shape $\times$ intent class (-0.0005) and OakInk2’s scene $\times$ verb (-0.0001) return +2.3% and +3.0% over forty seeds each ($p = 0.60$ and 0.96). Both are significant against zero ($p = 0.021$ and 0.015), which is why the reading is against the control: the criterion “distinguishable from zero” would have counted the instrument’s bias as difficulty. The pre-registration declared the test against zero and a point floor of 2.9%; the comparison of distributions was adopted after both results were seen, and we say so. Sufficiency is false: GRAB’s shape $\times$ fine intent has interaction at $z = 3.9$ and returns +0.7%, at the control.

The screen does not predict magnitude. We swept six real axes before using it to choose one, and over the five that survive withdrawal the interaction excess tracks the baseline arm’s penalty at $r = 0.843$, already not significant at that sample size ($p = 0.073$). We then swept the class axis on the screen’s own recommendation, at twelve seeds, with the reading declared in advance: ranked fifth of six, it predicts a small penalty. It returned +0.01650, the largest difficulty we measured anywhere, and that one sweep took the association to $r = 0.463$ ($p = 0.355$). The obvious escape does not hold: if the screen failed when the left factor is coarse, the id-

prefix axis would fail too, and it has the same five groups and is ranked correctly. We therefore use the screen to decide which axis to try first, and let the sweep say how hard it is. Fig. 2 gives the eight axes as measured quantities, difficulty beside advantage.

A withdrawn axis. DexYCB’s sweep is excluded because it could not have measured what it claimed. It ran at $\text{min_chains} = 1$, which disables the restriction, as all sixty transitions in its bundle fall in a single fine group. The coverage gate then fails on all twelve seeds at both budgets, the informed arm reaching a mean of 1.58 of the 4 held compositions. The failure is one of breadth, not depth, as the compositions it reaches carry 5.0 to 16.0 examples, inside the 3.8 to 37.3 band the other axes sit in. But a held composition the informed arm never sees is what Table I calls pinning the penalty at zero by construction, and no configuration recovers it. DexYCB’s null belongs to the split, and the dataset is untested here. A second sweep is withheld for the same reason and given with the release.

Four further non-positive-excess axes cannot be tested at all. The leak-free split of Section III-C holds out a transition only if it appears in several distinct fine labels; on DexYCB’s subject $\times$ object axis and on GRAB’s two object axes every transition appears in exactly one, so nothing can be held out at any setting, and GRAB’s object $\times$ intent class clears that bar only at a coverage the gate rejects on 28 to 33 of 40 rows. The three necessity tests above are therefore the three these datasets afford, not a sample of them.

Coverage. We report the remaining axes case by case. Our own gate transcripts for two of them covered ten seeds of sweeps that used forty and twelve, so the claim they supported stood for longer than the evidence did. GRAB’s shape axis and OakInk2’s pass every row of their full sweeps, 80 rows and 24 respectively, and the two post-submission axes pass 39 and 38 of 40; these carry the nulls the necessity argument rests on. Thinness costs a penalty its size, so the class axis is measured conservatively, failing 6 of 24 rows on breadth and still returning +0.01650. The other three fail a minority of rows whose informed arms still reach 4 of 5 held compositions at 8.0 to 28.0 examples each, against DexYCB’s 1 or 2 of 4. Excluding every seed flagged moves the confirmatory advantage from +0.00397 over 69 to +0.00419 over 66 ($t p = 0.0010$), and the release gives the rule with the per-row tallies.

What produces it. We do not know which collection property produces the interaction, and two candidates fail. Neither more pose-relevant intents nor a crossed design accounts for it, and the release refutes both on the datasets’ own numbers. At $n = 11$ the interval on a correlation is wide, so neither is refuted outright, but neither is what the data show. The axes returning a penalty all belong to one dataset we cannot distinguish from the others on any property we measured.

C. Primitive identity

Reconstruction error says nothing about whether the bank’s units are units. The identity ratio is between-primitive




Fig. 2. Left, how much the held-out compositions cost a baseline that never saw them. Right, how much of that cost the modular arm removes, paired by seed. Faded rows are axes whose difficulty is indistinguishable from the zero-truth control, where an architecture comparison is uninformative rather than negative. Every value is read from the released run artifacts.




Fig. 3. Four of the sixteen primitives in one bank, each decoded from the same start pose and shown at five points through its window, with the trailing fingertip paths faded. The last row is one primitive followed by another, a pairing the bank never saw in training.

over within-primitive variance in what a primitive produces (Fig. 3); above 1, the output is determined more by which primitive it is than by what preceded it. Pooled over 13 runs carrying Section III-B’s consistency term it is 1.841 ± 0.692 , with 11 of 13 above 1 ($t = 4.38$, $p = 0.0009$; sign test $p = 0.0225$) and a range of 0.87 to 2.84. Those runs are trained on OakInk2, which Section V-A finds null; four banks trained on OakInk-Image under an identical recipe give 1.848, all four above 1 ($t = 5.07$, $p = 0.015$), at an $n = 4$ where the sign test cannot clear 0.05. Read the ratio as a target being met, since the consistency term optimizes exactly this quantity.

The ten banks trained without the term differ from the treatment arm in size or dataset, so we do not argue from them. Two smaller sweeps that hold size and dataset fixed at

$K = 12$ and vary only the weight both move the same way, mean identity rising from 0.677 to 1.773 in one and from 0.299 to 1.110 in the other, over three seeds each.

D. Transfer to the Shadow Hand

Setup. We attempted twenty-two comparisons between motion generated for unseen primitive pairs and for seen ones; a pair is either seen or unseen, so the pairing of Section III-C is unavailable here. Grasp retention is scored by Fisher exact, the rest by Welch’s unpaired t -test. Five are undefined because both arms are identically zero. Of the remaining seventeen, ten on grasp retention, four on abduction cost and three on penetration, sixteen fail to separate the arms; the seventeenth separates them the wrong way, abduction cost $p = 0.037$ with the unseen arm worse, 3.32° against 4.64° . Four tests at $\alpha = 0.05$ give one such result about one time in five, consistent with no effect, and we carry it as a contradiction.

Failure to reject is not evidence of similarity, so we test equivalence directly, and it does not survive its own sensitivity check. Pooled over five priors at 165 clips per arm, abduction cost is equivalent within 1° ($p = 0.008$) against a margin fixed in advance. But 60 clips per arm come from one prior whose correction never fired, all 120 values exactly 0.0; that block cannot fail a one-degree margin and collapses the pooled variance, and dropping it leaves 105 clips per arm at $p = 0.124$. Dropping instead the prior that carries the reversal moves the difference from -0.36° to -0.06° , so two blocks carry the test in opposite directions; the release gives the power calculation. The other margins say little: penetration never exceeds 2.3% after correction, retention within ten points is not established at easy ($p = 0.069$) or medium ($p = 0.074$), and the pass at hard is vacuous at 3.3% against 0.0%.

The retargeting baseline. Direct retargeting of the first 40 trajectories per dataset interpenetrates in 24.6%, 76.8% and 59.0% of frames on OakInk-Image, DexYCB and GRAB; the abduction correction resolves 92%, 97% and 64%, leaving 2.1%, 2.1% and 21.2%. The failure concentrates by finger

pair, index against middle on OakInk-Image (1,617 of 3,960 penetrating contacts) and DexYCB, little against ring on GRAB. GRAB’s thumb also arrives most clipped, THJ4 saturating in 87% of frames because the robot thumb cannot extend past zero; that is joint travel the correction cannot reach.

At the task level, four modular banks against four monolithic priors, sampled unconditionally on a MuJoCo grasp-retention task, retain 28.1% and 27.7% respectively at a pre-declared 480 clips per arm ($p = 0.94$), against 25.0% for uniform noise. Those monolithic priors carry 0.59 to 0.63 active units against the 0.996 to 1.000 of Section III-B’s baselines, so the null is between a healthy bank and a partly collapsed latent. An exploratory run on the same eight models at 160 clips per arm gave 35.6% for the modular arm against 26.9% for the monolithic one (Fisher $p = 0.117$). It appears here because the confirmatory script declared it would, that test existing precisely because the p value had been seen.

VI. LIMITATIONS

Section V-A’s two metrics disagree, and the negative control disposes of both. Matching takes the raw effect from +0.00693 to +0.00254, past neither test and down to that arm’s floor. On the normalized penalty it appears to remove nothing, yet the same contrast on `pc_easy` returns +0.099. Normalizing was meant to correct for a baseline that reconstructs worse, and the rate-matched baseline reconstructs better, so the correction runs the wrong way. No arm holds rate and reconstruction quality fixed at once, and no metric here separates the effect from the instrument. Intervening on the modular side is worse still. Charging the assignment its categorical KL zeroes that channel and raises reconstruction error by 57 to 72%, so its -0.00336 for `perframe_rate` against `modular_rate` (sign test $p = 0.041$, interposable control) describes a model we broke.

Two sweeps of the same axis differ by more than either’s standard error, and this design separates neither cause. Training three arms consumes the random stream differently from training two, and the sweeps use `min_per_composition` 4 and 5, which changes the informed set and leaves the naive set identical. Read across sweeps, effect sizes carry that spread, whereas Section V-A’s rate-matched contrast sits within one sweep and is sound.

The paired penalty also returns a value where the truth is zero. `pc_easy` has no via-point offset, so its compositional penalty is zero by construction, yet in its own rate-matched bundle the baseline arm reports +0.00258 ($p = 0.0002$) at budget 256 and +0.00306 ($p < 0.0001$) at budget 64 over twenty seeds. Taking the budget-256 floor throughout, that is 25% of the unmatched baseline’s difficulty in the re-sweep (+0.01024, where Table II gives +0.01042 for the original) and 44% of the rate-matched baseline’s (+0.00586). Difficulty therefore reads partly off the measuring model, which is why Section V-B reads it against this control and not against zero. The same control bounds the advantage, and on one arm it bounds it away. Against the rate-matched

baseline the raw floor is +0.00155, 61% of the whole effect, and the normalized floor is +0.099 at 18 of 20, which clears both declared tests. No headline should be read without these.

We truncated one sweep and extended another after seeing its p value. The category sweep declared 70 seeds, 69 of which returned both arms, and stopped at 52. The missing 17 are null alone (-0.00046, $t p = 0.86$, 11 of 17 still favoring modular), so the first 52 were the favorable half and pooling moves d from 0.60 to 0.41. The affordance axis has the opposite defect. Seventy seeds gave +0.00270 at sign $p = 0.120$, sixty more were run, and the pooled +0.00365 passes at sign $p = 0.0003$. Having looked survives completing the sweep, so that axis is exploratory and the confirmatory claim stays with the category axis. The effect is small per split, too, as the modular arm wins 48 of 69 seeds on the category axis and 86 of 130 on the affordance axis, failing on roughly three splits in ten.

Transfer is kinematic throughout: contacts are read off after `mj_forward` places the model, and nothing measures forces. Its baseline is weak in a specific way. Since DexPilot [19] the retargeting work we are aware of operates in task space, and the widely used dex-retargeting implementation of that objective, released with AnyTeleop [21], carries no collision term. The 24–77% interpenetration figure therefore characterizes that baseline and not the field’s standard, so we make no claim about pipelines carrying a task-space objective. We ran one, in which both arms penetrate in 100% of frames over four clips each, undefined for the reason Section V-D gives for its other zero-versus-zero cells. We report it because it cuts our way and cannot support us. Those banks were trained on OakInk2, which Section V-A finds null.

Five further limitations are documented with the release, and two of them bound this paper’s claims directly. The advantage does not grow as data shrinks, over four sweeps at budgets 64 and 256 with two reversing and none significant. And contact does not explain the nulls, as GRAB’s 16 per-region channels leave the penalty at +0.00032 ($p = 0.813$).

The experiment that would settle this was not run. Sampling both priors unconditionally gives the architectures no channel through which to differ on grasp retention, which is why Section V-D’s task-level comparison cannot see the effect. Seeing it needs each prior used as the action space for a policy that is trying to grasp, as LAMP [1] and DexSyn-Refine [7] both do for continuous latents. LAMP’s discrete VQ-VAE arm finds gesture jumps at codebook switches, but it quantizes the commanded pose where our bank keeps a continuous decoder and makes only the selection discrete, so that artifact is not entailed. A bank of this form against a parameter-matched continuous latent is what we would run next.

VII. CONCLUSION

Before spending GPU on a compositional axis, run the interaction statistic of Section III-D over the raw pose arrays. It trains nothing and finishes in minutes, and a non-positive excess is a reason to look elsewhere: on eight axes across

three datasets no axis carried difficulty without positive interaction, and the three that could have shown otherwise sit at the zero-truth control. A positive excess licenses less, as GRAB’s shape axis clears the same bar and sits at the control too. Read every penalty against that control; at forty seeds a test against zero detects the instrument, not the composition. Where difficulty is present the modular bank pays the smaller penalty, but at matched rate we claim nothing for modular structure, since our own control manufactures most of the effect we would be claiming. A reader who wants that settled should look to a bank used as a policy’s action space, where the architectures were proposed and where we did not test them. We stop short of a dataset design rule. Every axis that returned a penalty belongs to one dataset, and no collection property we measured predicts which.

Acknowledgment

Claude (Anthropic) was used throughout this work under the authors’ direction: it drafted and revised the manuscript text, wrote the analysis, gating and build scripts in the accompanying repository, and launched the experiment runs whose artifacts the paper reports. The authors designed the study, chose every axis, gate and pre-registered test, reviewed every retained sentence and number against the run artifacts, and take full responsibility for the content. The two sweeps of Section V-B marked as post-submission were run after the initial submission under the protocol pre-registered with the release. Code, run artifacts and an extended appendix will be released with the published version.

REFERENCES

- [1] X. Yang et al., “LAMP: Latent Motion Prior-Guided Real-World Learning for Dexterous Hand Manipulation,” arXiv:2607.06323, 2026.
- [2] T. Wang et al., “MotionBricks: Scalable Real-Time Motions with Modular Latent Generative Model and Smart Primitives,” ACM TOG / SIGGRAPH 2026, arXiv:2604.24833.
- [3] S. Linderman, M. Johnson, A. Miller, R. Adams, D. Blei, L. Paninski, “Bayesian Learning and Inference in Recurrent Switching Linear Dynamical Systems,” AISTATS 2017, PMLR 54:914–922.
- [4] Z. Dong, B. A. Seybold, K. Murphy, H. H. Bui, “Collapsed Amortized Variational Inference for Switching Nonlinear Dynamical Systems,” ICML 2020, PMLR 119:2638–2647.
- [5] H. Zhang, S. Starke, T. Komura, J. Saito, “Mode-Adaptive Neural Networks for Quadruped Motion Control,” ACM TOG 37(4), Article 145, SIGGRAPH 2018.
- [6] O. Taheri et al., “GRAB: A Dataset of Whole-Body Human Grasping of Objects,” ECCV 2020, arXiv:2008.11200.
- [7] H. Lee, H. Jung, S.-H. Heo, S. Yang, “DexSynRefine: Synthesizing and Refining Human-Object Interaction Motion for Physically Feasible Dexterous Robot Actions,” arXiv:2605.05925, 2026.
- [8] N. B. Gutierrez, J. M. Cloud, W. J. Beksi, “Movement Primitives in Robotics: A Comprehensive Survey,” arXiv:2601.02379, 2026.
- [9] B. Jiang, W. P. Reger, R. Martin-Martin, “CoDex: Learning Compositional Dexterous Functional Manipulation without Demonstrations,” arXiv:2606.31909, 2026.
- [10] H. Zhang, J. Ferchow, J. Song, M. Meboldt, “UniCross: Unified Cross-Skill Dexterous Manipulation Synthesis,” arXiv:2607.28198, 2026.
- [11] A. E. Tekden, Y. Bekiroglu, “Compositional Motion Generation from Demonstration with Object-Centric Neural Fields,” arXiv:2607.07129, 2026.
- [12] J. M. Lu, R. Zhang, I. Lee, D. Alvarez-Melis, E. Pavlick, N. Saphra, “Adversarial Concept Search: Predicting Compositional Errors From Feature Geometry,” arXiv:2606.13934, 2026.
- [13] C. Pan et al., “SPIDER: Scalable Physics-Informed Dexterous Retargeting,” arXiv:2511.09484, 2025.
- [14] L. Yang et al., “OakInk: A Large-scale Knowledge Repository for Understanding Hand-Object Interaction,” CVPR 2022, arXiv:2203.15709.
- [15] X. Zhan et al., “OakInk2: A Dataset of Bimanual Hands-Object Manipulation in Complex Task Completion,” CVPR 2024, arXiv:2403.19417.
- [16] Y.-W. Chao et al., “DexYCB: A Benchmark for Capturing Hand Grasping of Objects,” CVPR 2021, arXiv:2104.04631.
- [17] E. Todorov, T. Erez, Y. Tassa, “MuJoCo: A Physics Engine for Model-Based Control,” IROS 2012, pp. 5026–5033.
- [18] K. Zakka, Y. Tassa, and MuJoCo Menagerie Contributors, “MuJoCo Menagerie: A Collection of High-Quality Simulation Models for MuJoCo,” 2022. The Shadow Hand E3M5 model used here is the repository’s shadow_hand entry. https://github.com/google-deepmind/mujoco_menagerie
- [19] A. Handa et al., “DexPilot: Vision-Based Teleoperation of Dexterous Robotic Hand-Arm System,” ICRA 2020, pp. 9164–9170, arXiv:1910.03135.
- [20] J. Romero, D. Tzionas, M. J. Black, “Embodied Hands: Modeling and Capturing Hands and Bodies Together,” ACM TOG 36(6), Article 245, SIGGRAPH Asia 2017.
- [21] Y. Qin et al., “AnyTeleop: A General Vision-Based Dexterous Robot Arm-Hand Teleoperation System,” RSS 2023, arXiv:2307.04577.
- [22] E. Jang, S. Gu, B. Poole, “Categorical Reparameterization with Gumbel-Softmax,” ICLR 2017, arXiv:1611.01144.
- [23] I. Higgins et al., “beta-VAE: Learning Basic Visual Concepts with a Constrained Variational Framework,” ICLR 2017.