

Joint-centre offset dominates 3D position error and can reverse comparisons between markerless triangulation methods

Title page

[Author names — given name(s) and family name(s) of each author, in the order used in the submission system. To be completed.]

[Affiliations — full name and postal address of each affiliation, including the country, keyed to the author names by superscript letter. To be completed.]

[Corresponding author — name and email address. The Guide for Authors requires the email on the title page, and it is published with the article. To be completed.]

Abstract

Markerless motion capture pipelines are commonly tuned and compared using three-dimensional joint position error measured against a marker-based reference. However, pose-estimation keypoints and marker-based joint centres are not the same anatomical construct, and at the hip and knee the two have been reported to differ by several centimetres, so it is unclear what such a comparison actually measures. We held a single pipeline fixed and varied four triangulation decisions (confidence weighting, rejection by reprojection error, rejection by object-space error, and Gauss–Newton refinement) together with a population offset correction, one at a time, across 104 overground walking trials from 11 participants of the BioCV dataset, recorded on a nine-camera ring against a synchronised 15-camera marker-based reference. At the hip and knee the keypoint lay 27–37 mm from the reference joint centre, which accounted for 74–81% of squared position error within a participant and for 47–70% when pooled across participants, compared with 2.5 mm for the largest triangulation effect measured. When each participant's own offset was subtracted, per-joint verdicts changed. Both rejection criteria appeared to improve hip position as measured, but worsened it by 1.7 and 2.0 mm once the offset was removed, leaving the estimate noisier about its own mean yet better centred on the reference. No triangulation effect at the hip, knee or ankle exceeded 2.5 mm and no single-step angle effect exceeded 0.38° on the four chains tested. All contrasts came from a single rig, task and laboratory, and the family structure was exploratory. These findings suggest that per-joint position verdicts should be interpreted in the light of the joint-centre offset they are likely to contain.

Keywords: markerless motion capture; multi-view triangulation; gait analysis; joint centre; validation methodology

I. INTRODUCTION

Markerless motion capture has been reported to agree with marker-based systems to within a few degrees for lower-limb kinematics (Pagnon et al., 2022; Uhlich et al., 2023), although this evidence is concentrated in the sagittal plane (Cheng et al., 2025) and agreement is considerably poorer for rotations about the long axis of a segment, which Wade et al. (2022) place at 6.9–13.2° (see also Varcin and Boocock, 2026). As these systems move towards clinical and field use, the decisions made inside the reconstruction pipeline, such as how camera views are combined and whether unreliable views are discarded, are increasingly being tuned and compared.

The metric used for that tuning has largely been inherited from computer vision. The benchmark that established the convention reports angle error alongside position error (Ionescu et al., 2014), but pipeline decisions are typically checked against position alone. Of four studies that vary such a decision, two evaluate it on position while reporting angles (Karashchuk et al., 2021; Uhlich et al., 2023), one reports angles alone (Pagnon et al., 2021) and one position alone (Needham et al., 2021), and none relates the two. Since landmark misplacement propagates to joint kinematics unevenly across joints and planes (Della Croce et al., 2005), an intervention that displaces an angle's three defining joints by different amounts need not change that angle proportionately.

A further difficulty is that pose-estimation keypoints and marker-based joint centres are not the same anatomical construct. Needham et al. (2021) reported differences of 30–50 mm between the two at the hip and knee, and Kanko et al. (2021) independently observed 3.6 cm at the hip against 2.2–2.4 cm at the knee and ankle, with differences that varied over the gait cycle. An offset of this kind is not produced by triangulation and is common to any two pipeline variants being compared. Whether a difference in position error between such variants may therefore be dominated by the shared offset, and not by the decision under test, does not appear to have been examined.

The aim of this study was therefore to quantify how much of the position error at the hip and knee the keypoint-to-joint-centre offset accounts for, and whether comparisons between triangulation variants change when that offset is removed. We held a single pipeline fixed and varied five components one at a time: confidence weighting, rejection by reprojection error, rejection by object-space error (Lu et al., 2000), Gauss–Newton refinement, and a population offset correction. We also asked whether a gain in position transmits to joint angle, testing each claimed dissociation as an interaction, since two separate significance verdicts do not establish one (Gelman and Stern, 2006).

II. METHODS

A. Dataset

The BioCV dataset was analysed (University of Bath Research Data Archive, doi:10.15125/BATH-01258; Evans et al., 2024). Eleven of the fifteen healthy volunteers the archive describes were present in the copy released to us under a data use agreement. These eleven (six male, five female, aged 26.2 ± 5.1 years) each performed ten straight-line overground walking trials at a self-selected pace, except one who performed nine. Five trials were lost to missing reference data or unusable video, leaving 104, from each of which 60 near-evenly spaced frames were retained (6,240 frames). Because retained frames were 45 ms apart, no temporal filter could be applied.

Video was recorded by a nine-camera ring at 1920×1080 and 200 Hz. The lenses distort strongly, and correction under the model of Brown (1971) displaces a point by 6.09 px on average and 84.80 px at worst, so all results are conditional on it (Table S2b). Reference kinematics came from a synchronised 15-camera Qualisys Oqus system recording at 200 Hz and processed with the dataset's own six-degree-of-freedom model. The two systems were hardware-triggered and a whole-frame offset was corrected post hoc from an LED pulse. The providers do not quantify the

remaining sub-frame residual, which we bound at typically 3.6 mm at the hip and knee, and their reported cross-system alignment error of 3.62 mm was not recomputed (Tables S12, S14).

Keypoints were detected with RTMPose (Jiang et al., 2023) and each COCO-17 keypoint was differenced against the identically named output of the reference model, so every offset reported here is the keypoint minus the model joint centre. That model constructs joint centres as the midpoint of the medial and lateral markers at every joint except the hip, where it applies the regression of Bell et al. (1989). Two further detectors, RTMO and HALPE-26, were run on the same frames to test whether the offset was specific to a single model.

B. Pipeline arms

An arm is a complete pipeline configuration. All arms shared detections, calibration and candidate cameras, differing only in the named component, and contrasts are named with the baseline first. We varied confidence weighting (the detector's confidence scores against uniform weights), the two rejection criteria described below, Gauss–Newton refinement against the direct linear transform, and a population offset table, a fixed per-joint translation fitted to this rig and held out in participant only. The distance weightings and matched-retention arms use ground-truth quantities and are diagnostic controls, so claims about a deployable rule rest on the adaptive comparison alone (Table S10). Heading varied by 1.0 – 14.1° within ten of the eleven participants, and the world- and body-frame decompositions differ by at most one percentage point in the lower limb (Tables S3m, S20). The two estimators minimise Σwr^2 and $\Sigma w^2 r^2$ respectively, and contrasts that hold this fixed are labelled w^2 -matched.

C. The two rejection criteria

For camera i the reprojection residual is the norm of the difference between the projected point and the detection, in pixels, whereas the object-space residual is the perpendicular distance in millimetres from the reconstructed point X to the ray through that detection, which is the collinearity error minimised by Lu et al. (2000). Those authors report their algorithm to be more robust to outliers than Levenberg–Marquardt alternatives, which is a claim about an estimator and not about either residual's power to identify a bad observation, so the comparison drawn here is our own. They also note that this error grows with distance from the camera, and for small angular residuals

$$e^{obj} \approx (Z/f) \cdot e^{px} \quad (1)$$

where Z is the depth of X from that camera and f is its focal length in pixels. Object-space error is therefore a depth-scaled reprojection error, and on a given rig the two criteria differ only through the way Z/f is distributed across the cameras.

Both criteria used the same adaptive rule. The camera with the largest residual was discarded whenever that residual exceeded the median plus three scaled median absolute deviations, the point was re-triangulated after each discard, and the procedure stopped at a floor of four cameras. In the matched-retention arms this rule was replaced by a fixed count, so only the ordering of discards differed. Over a multiplier swept from 2.5 to 4.0 the thigh–shank angle cost varied from 0.053° to 0.286° , so its direction is robust to the threshold and its magnitude is not (Table S21).

D. Outcomes

Position error was measured on two joint sets: a peripheral set of both ankles and both wrists, labelled [pos:AW] in the tables, and a lower-limb set of the hip, knee and ankle of each side, labelled [pos:HKA]. The latter was accumulated inside the angle computation so that it shared frames and arms with the angle outcomes.

The angle outcomes were three-point included angles and no kinematic model was fitted. They were computed identically for the estimate and the reference and named for the segments defining them: the thigh–shank angle, angle(hip, knee, ankle), which approximates knee flexion; its frontal-plane form; and the trunk–thigh angle, angle(mid-shoulder, hip, knee), which approximates hip flexion but also moves with trunk lean. Four baselines for the thigh–shank angle circulate here on different arms and frame sets, reconciled in Table S16, and none is directly comparable in absolute value with model-based reporting. The lower-limb position set contains all three points of the thigh–shank angle but only some of those of the other two, so we claim a match between a position outcome and an angle's defining joints for the thigh–shank angle only. Four chains are computable under COCO-17, the three angles above and the elbow, and where a chain's terminal keypoint admits more than one choice a pre-declared veto withholds the verdict unless both choices agree in sign (Table S13).

E. Statistics

The participant was the inferential unit. Each contrast was reduced to one mean difference per participant and all $2^{11} = 2048$ sign patterns were enumerated, giving an exact two-sided p with a floor of 0.000977 that certifies only that all eleven participants agreed in direction. Multiplicity was controlled with the Benjamini–Hochberg procedure

(Benjamini and Hochberg, 1995) within each of fourteen declared families (Table S4b), under the positive regression dependence we assume without demonstrating (Benjamini and Yekutieli, 2001). Under the arbitrary-dependence modification, which we did not adopt, 39 verdicts fall and are marked † wherever quoted (Table S5). The families were not pre-specified but grew as claims required contrasts, each being re-run in full, so this study is exploratory in its family structure.

Because significance in one outcome and not in another does not by itself establish a dissociation (Gelman and Stern, 2006), each was tested as an interaction: one relative effect minus the other, computed per participant. The null of that test is equal fractional change, so differences are also reported in millimetres and degrees, and an effect surviving on only one scale is scale-dependent.

The exact test assumes the two arms are exchangeable within a participant, which the shared frames and detections provide for the triangulation arms but not for the population offset table, whose correction is built from the other ten participants. A rule declared beforehand therefore withholds the q value on every leave-one-out row, leaving a nested null as their only inferential statement (one-sided $p = 0.026$). Every affected family was corrected again without those rows, and 143 of the 244 declared contrasts carrying a verdict survived, out of 247 declared in all (Tables S3e, S4b, S22). Under leave-one-participant-out refitting the fourteen headline contrasts each retained their direction at $p < 0.05$ (worst $p = 0.029$; Table S6). Intervals invert the sign-flip test, at 95% for a quoted magnitude and 90% for a bounded null (Tables S3k, S8).

III. RESULTS

A. The joint-centre offset dominates position error at the hip and knee

Agreement with the reference was 2.575° , 2.637° and 3.651° of mean absolute error on the three angles, and 28.5 mm on the lower-limb position set (Table S2d). Most of what that metric measured at the hip and knee appears not to be attributable to triangulation. The keypoint lay 27–37 mm from the reference joint centre (Table 3). Within a participant that offset accounted for 74–81% of squared position error at the hip and knee and 44% at the ankle. Pooled across participants, as a cohort-level validation study measures it, the between-participant scatter moves into the random term and the share falls to 47–70%. For comparison, the largest triangulation effect we measured at any of these joints was 2.5 mm, an order of magnitude smaller. A difference of norms does not remove a term common to both arms, and the

first-order account of what it leaves behind fails at the hip and ankle (Table S7), so the reversals reported below are an empirical result and were not anticipated.

Table 3. The joint-centre offset, and why correcting it fails
(a) Correcting the offset (mean absolute error, degrees)

angle	uncorrected	population (leave-one-out) correction	own-offset ceiling
thigh–shank angle	2.575	3.201	2.495
frontal-plane thigh–shank	2.637	2.768	2.486
trunk–thigh angle	3.651	3.363	2.541

The deployable correction worsens the thigh–shank angle by 0.606°, on a row whose q is withheld (§). The own-offset column is an upper bound requiring participant-specific marker data, and buys only 0.080° at the thigh–shank angle against 1.110° at the trunk–thigh angle.

The column arithmetic does not equal the tested effect, and should not. Differencing this table gives 0.626° at the thigh–shank angle, 0.131° at its frontal-plane form and 0.288° at the trunk–thigh angle, against tested effects of −0.606°, −0.149° and +0.164°. The table reports **means pooled over all observations**; the tested effect is the **mean of the per-participant differences**, which is the inferential unit throughout (Section 2.5). The two coincide only when every participant contributes equally, and participants here contribute between 6 and 10 usable trials (P26 lost four; P04 was archived with nine and P28 lost one trial to missing reference data). Unequal weighting accounts for the gap to

within rounding: computing the same three contrasts by pooling observations reproduces the column arithmetic (−0.626, −0.130, +0.288°), and computing them per participant reproduces the tested effects (−0.606, −0.149, +0.164°). The own-offset column reconciles closely at the hip, 1.110° pooled against 1.094° tested, and not at the thigh–shank angle, 0.080° against 0.166°, a factor of two. Where the two differ we quote the per-participant value, since that is what was tested.

(a2) How much of the position error is a systematic offset — bias/variance split of the 3D error vector, world frame, on the adaptive object-space-rejection arm with confidence weights. The offset is estimated WITHIN each participant and then averaged across them, so between-participant variation in the offset does not enter the random term; the body-frame split of Table S3m differs by less than 0.5 mm

joint	RMSE (mm)	systematic offset (mm)	random residual (mm)	offset share of squared error
hip	41.2–41.3	35.4–36.9	18.6–21.1	74–80%
knee	31.0–33.4	27.5–30.1	14.3–14.6	79–81%
ankle	17.3–17.4	11.5	13.0	44%

Pooled across participants, which is what a cohort-level validation study measures, the between-participant scatter of the offset (hip 20.0–21.9 mm, knee 11.2–13.4 mm, ankle 6.2–6.7 mm) moves out of the offset term and into the random one. The share is then **47–57%** at the hip, **61–70%** at the knee and **30–32%** at the ankle, in the world frame declared above. Both columns are reported because they answer different questions: this paper's contrasts are differences within a participant, where that participant's own offset is common to both arms, so the within-participant share is the one they rest on.

component	mean magnitude (mm)	effect on flexion
axial, along the limb	27.45	changes limb length, not flexion
in-plane perpendicular	11.04	drives flexion error
out-of-plane	8.63	affects ab/adduction, not flexion
total loffsetl	32.36	—

Each entry is a mean of per-frame absolute values, so the components do not combine into the total by root sum of

Every range in this panel is the span between the left and right sides, not a confidence interval. The per-joint effects contrasted in Table 1(b) span 0.024 to 1.564 mm. Decomposing the estimator change as well gives 2.415 mm at the hip for DLT to Gauss-Newton and 2.469 mm for adding rejection under it (Table S8), so the largest per-joint triangulation effect at the hip, knee and ankle is 2.5 mm, one order of magnitude below the offset at hip and knee.

(b) Why the thigh–shank angle gains so little — knee offset decomposed relative to the hip–ankle chord

squares; the total is the mean of the per-frame offset magnitude. That is also why 32.36 mm exceeds the 27.5–

30.1 mm systematic knee offset of Table 3(a2), which is the norm of the mean: a mean of magnitudes carries the per-frame scatter that a norm of a mean removes. The two tables do not disagree.

joint	loffsetl across quintiles (mm)	change	<i>q</i>
hip	36.4 → 38.6 → 38.7 → 42.9 → 43.0	+6.19	0.0065
knee	32.2 → 32.3 → 29.1 → 28.1 → 28.5	-2.82	0.029†
ankle	19.1 → 16.3 → 13.6 → 14.0 → 14.6	-4.27	0.013

The quintile columns are means pooled over observations; the change column is the mean of the per-participant differences, which is the tested quantity (§2.5). The two differ because participants contribute unequal numbers of trials, exactly as set out under Table 3(a).

Individual components move further than the magnitudes, because direction rotates as well as length: the hip's anterior component by -13.60 mm ($p = 0.001$) and the knee's superior component by -8.97 mm ($p = 0.001$).

RTMO and HALPE-26 reproduced the offset on the same frames, and all three detectors disagreed with the reference in the same direction for every participant, although they retain COCO-17 landmark semantics and so are not independent replications (Tables S9, S13). At the hip the mediolateral component of the offset was bilaterally symmetric and directed outward (-18.9 and +18.5 mm), a pattern the Bell et al. (1989) regression used by the reference would be expected to produce, whereas a shared timing lag or common-mode registration error would not.

B. Removing the offset changes per-joint verdicts

Subtracting each participant's own offset at each joint moved verdicts in both directions (Table 4). Both rejection contrasts reversed at the hip, the object-space criterion moving from +0.608 to -1.726 mm and the reprojection criterion from +0.633 to -1.983 mm. Because pairing two separate verdicts would not establish a reversal, the difference itself was tested, giving -2.335 mm on the object-space criterion and -2.616 mm on the reprojection criterion per participant, both at the enumeration floor and surviving correction within a declared family of eighteen ($q = 0.003$).

(c) **The offset varies with posture** — binned by quintile of the ground-truth thigh-shank angle; highest minus lowest quintile, per participant, exact sign-flip

The as-measured halves of those pairs are uncorrected ($p = 0.023$ and 0.044), so the reversal rests on the difference. Two further contrasts reversed at the knee, at matched retention $k = 1$ and $k = 2$.

Table 4. The per-joint contrasts repeated on offset-removed error

Each participant's own mean error vector for a joint is subtracted from both arms of a contrast, so what remains is the perturbation and not the displaced target. Subtracting one arm's own centroid centres that arm and not the other, which biases the contrast downward when the offset comes from the first-named arm and upward when it comes from the second. The midpoint of the two is symmetric and carries every verdict; the single-arm range is the span of that bias and is shown for reference only — a range must not be built from it. *As measured* repeats Table 1(b) for the first four contrasts; the last two are the rejection contrasts the Conclusion's position null rests on, added under a rule fixed before they were run.

Effects are millimetres, positive meaning the second-named arm is better. p is the exact sign-flip test over all 2^{11} sign patterns; q is Benjamini-Hochberg within $m = 18$, run once per reference column so that each column is a complete correction over the same eighteen contrasts. Bold marks a midpoint effect surviving that correction, and † one that does not survive the arbitrary-dependence sensitivity of Table S5 — the same convention as Tables 1 and 2, applied here so the absence of daggers is not read as a property of the panel. No row has a midpoint verdict differing from both single-arm verdicts. The component decomposition, the single-arm columns in full, the per-subset control and the dispersion/centroid split are in Table S7.

contrast	joint	as measured	<i>p</i>	offset-removed	<i>p</i>	<i>q</i>	single-arm range
adaptive criterion	hip	-0.024	0.8350	+0.264	0.0010	0.0035	[+0.221, +0.308]
adaptive criterion	knee	+0.339	0.0020	+0.181	0.0547	0.1094	[+0.170, +0.192]
adaptive criterion	ankle	+0.190	0.0352	+0.150	0.1084	0.1774	[+0.142, +0.158]
criterion at matched <i>k</i> =1	hip	+0.253	0.0840	-0.718	0.0010	0.0035	[-0.800, -0.632]
criterion at matched <i>k</i> =1	knee	-0.322	0.0029	+0.399	0.0039	0.0100	[+0.372, +0.426]
criterion at matched <i>k</i> =1	ankle	+0.134	0.1484	+0.076	0.3066	0.3943	[+0.063, +0.088]
criterion at matched <i>k</i> =2	hip	+0.530	0.0020	-0.108	0.5771	0.6790	[-0.222, +0.009]
criterion at matched <i>k</i> =2	knee	-0.655	0.0020	+0.738	0.0029	0.0088	[+0.643, +0.828]
criterion at matched <i>k</i> =2	ankle	+0.219	0.0762	+0.215	0.0664	0.1195	[+0.181, +0.247]
uniform → confidence	hip	+1.564	0.0010	+0.253†	0.0088	0.0198	[+0.069, +0.435]
uniform → confidence	knee	+0.365	0.0010	-0.024	0.7266	0.7693	[-0.031, -0.017]
uniform → confidence	ankle	+0.822	0.0010	+0.912	0.0010	0.0035	[+0.903, +0.921]
no rejection → object-space rejection	hip	+0.608	0.0234	-1.726	0.0010	0.0035	[-1.878, -1.564]
no rejection → object-space rejection	knee	-0.546	0.0039	-0.020	0.8486	0.8486	[-0.086, +0.046]
no rejection → object-space rejection	ankle	+0.426	0.0039	+0.256	0.1621	0.2432	[+0.213, +0.299]
no rejection → reprojection rejection	hip	+0.633	0.0439	-1.983	0.0010	0.0035	[-2.240, -1.714]
no rejection → reprojection rejection	knee	-0.885	0.0010	-0.203	0.2051	0.2840	[-0.279, -0.125]
no rejection → reprojection rejection	ankle	+0.237	0.1738	+0.107	0.6035	0.6790	[+0.084, +0.130]

What the table shows. At the hip the adaptive criterion and both rejection contrasts change sign once the target is put back, and the two rejection reversals are the ones Section 3.2 reports; their as-measured values are uncorrected and outside the declared per-joint family. At the knee the two matched-*k* contrasts reverse with both directions surviving correction. The ankle moves nowhere.

Decomposed by joint, the adaptive criterion improved the knee by 0.339 mm ($q = 0.004$) and seven of twelve per-joint contrasts survived correction (Fig. 2, Table S3b). On both

deployable rejection contrasts the hip gained and the knee lost while the three-joint mean was null, so opposed per-joint effects cancelled in the mean, although under correction this held only on a matched-retention control. One intervention also survived a comparison between joint subsets, on both scales and under arbitrary dependence. Confidence weighting differed by 1.05 percentage points and 0.276 mm between {hip, knee, ankle} and {hip, ankle} ($q = 0.004$; Table S3n), so the joint set a study summarises over changes what it reports.

A null in the three-joint mean can be two significant effects of opposite sign




Fig. 2. On a retention control, a null in a three-joint mean is two significant effects of opposite sign. 3D position effect at each lower-limb joint separately, with the three-joint mean at right. At matched retention $k = 2$ (left) the object-space criterion improves the hip and degrades the knee, both $q = 0.004$, cancelling to a null mean; under adaptive rejection (centre) the same criterion helps the knee (0.339 mm, $q = 0.004$), and the ankle effect (0.190 mm) is significant unadjusted but not after correction ($q = 0.053$). Confidence weighting (right) improves every joint. Bars are error against the reference as

measured; filled bars survive Benjamini–Hochberg correction within their family and open bars do not, which is why the ankle bar is open at an unadjusted p of 0.035. What happens once the offset is removed is in Section 3.1 and Table 4.

C. Camera rejection

Against no rejection, adaptive rejection left lower-limb position essentially unchanged, at 0.005 mm worse for reprojection and 0.163 mm better for object-space, bounded within ± 0.35 mm at 90%, but that bound was measured on a displaced target. With each participant's own offset subtracted, both criteria worsened hip error about that participant's own mean, by 1.726 mm for the object-space criterion and 1.983 mm for the reprojection criterion ($q =$

0.004). Scatter and centring are separate contrasts running in opposite directions. Scatter worsened by 1.801 mm for the object-space criterion and 2.082 mm for the reprojection criterion, while the centre of the estimate moved 1.648 and 1.958 mm closer to the reference, in that same order. Removing the offset within quintiles of the reference angle gave 1.924 and 2.235 mm in that same order, so the degradation is not the posture-dependent residual of Fig. 4. The knee and ankle remained null.




Fig. 4. The joint-centre offset is not fixed but varies with posture. Offset magnitude by quintile of the reference thigh–shank angle (180° = full extension for the three-point angle used here). Highest minus lowest quintile: hip +6.19 mm ($q = 0.0065$), knee -2.82 mm ($q = 0.029†$), ankle -4.27 mm ($q = 0.013$), each corrected within the thirteen-test posture family; † marks the one that does not survive the arbitrary-dependence sensitivity of Table S5.

The same rejection degraded the thigh–shank angle by 0.188°† and 0.141° and the trunk–thigh angle by 0.192°† and 0.167°†, reprojection first in each pair, which reverses the ordering the corresponding position pairs appear to take ($q =$

0.011–0.028; Table 2). On the matched-retention control the cost showed a dose-response reaching 0.747° and 0.798° at $k = 3$ (Fig. 3a).




Fig. 3. (a) Thigh–shank angle error against the number of cameras discarded per joint at matched retention, tested per participant against $k = 0$; all six steps survive correction within the dose-response family. Matched retention is a control: it fixes the number discarded at levels the adaptive rule never reaches, so the rate carries no deployable claim. The panel quotes the tested mean rate, not the marginal increment of Table 2(b). (b) The object-space advantage in peripheral-set position by camera-distance asymmetry, before and after lens-distortion correction (§3.2); the same control on the lower-limb set is in Table S2d. Filled bars have bootstrap intervals excluding zero.

Table 2. Joint angles

(a) **Contrasts** (degrees; positive = the second-named arm has lower error)

contrast	dep.	thigh–shank angle	frontal-plane thigh–shank	trunk–thigh angle
uniform vs confidence weighting	y	+0.023, n.s.	+0.078, $q = 0.004$	+0.227, $q = 0.006$
no rejection vs reprojection rejection, adaptive	y	–0.188, $q = 0.021^\dagger$	–0.013, n.s.	–0.192, $q = 0.022^\dagger$
no rejection vs object-space rejection, adaptive	y	–0.141, $q = 0.011$	–0.067, raw only	–0.167, $q = 0.028^\dagger$
reprojection vs object-space criterion, adaptive	y	+0.047, n.s.	–0.054, n.s.	+0.025, n.s.
confidence vs confidence $\times 1/d$ weighting, no rejection	n	+0.052, n.s.	–0.188, $q = 0.004$	–0.219, $q = 0.004$
reprojection vs object-space criterion, matched $k = 1$	n	–0.089, $q = 0.016^\dagger$	—	—
reprojection vs object-space criterion, matched $k = 2$	n	–0.129, $q = 0.013^\dagger$	—	—
no rejection vs reprojection rejection, under Gauss–Newton	y	–0.169, $q = 0.028^\dagger$	—	—
uncorrected vs leave-one-participant-out offset correction	y* ‡	–0.606, ‡ withheld	–0.149, ‡ withheld	+0.164, ‡ withheld

dep. marks an arm a validation study could deploy; **n** marks a diagnostic or control arm — the distance weightings use ground-truth distances, and matched- k fixes retention at a level the adaptive rule never reaches. **y*** marks the population offset table, which is deployable in principle but is computed here as a world-frame translation on trials whose headings span 56° across the study and 1.0 – 14.1° within ten of the eleven participants, so this particular table does not transfer to another laboratory or protocol (§2.3); its

dissociation is demonstrated within this rig. † marks a verdict that does not survive the arbitrary-dependence sensitivity of Table S5, i.e. one resting on the positive-dependence assumption. Exact 95% intervals for every surviving row are in Table S8.

(b) **Dose-response, thigh–shank angle** — mean absolute error pooled over observations (degrees), and the same steps tested per participant against $k = 0$, BH-corrected within a family of 6

cameras discarded per joint	0	1	2	3
reprojection criterion, pooled MAE	2.570	2.880	3.090	3.358
object-space criterion, pooled MAE	2.570	2.966	3.217	3.412
reprojection, tested cost (q)	—	+0.286 (0.011)	+0.482 (0.006)	+0.747 (0.006)
object-space, tested cost (q)	—	+0.375 (0.006)	+0.612 (0.006)	+0.798 (0.006)

The $k = 0$ baseline here is 2.570° where Table 3(a) and Table S2c give 2.575° for the same arm. The dose-response family is accumulated over the 12,468 observations on which every retention level is defined, a subset of the frames the main family uses, so the two baselines are the same quantity on different frame sets and are not reconciled to a single value.

All six steps survive correction. The tested cost averages **0.24–0.38° per camera discarded**, that is the cumulative cost at k divided by k ; the per-participant $k = 0$ baseline is

2.669° , against the 2.570° pooled mean above, for the reason set out under Table 3(a). Only the tested values are quoted in the text. Differencing the pooled row gives a MARGINAL increment of $0.20\text{--}0.40^\circ$ for each further camera, which is a different quantity from the mean rate above and is arithmetic on means rather than the contrast that was tested.

(c) **Dissociation tested as an interaction** — per participant, relative position effect minus relative angle effect, on one dimensionless scale; BH-corrected within the family of 20

intervention	dep.	Δ position	Δ angle	interaction	q
criterion at matched $k = 1$, thigh–shank angle	n	+0.07%	–3.21%	+3.29	0.005
confidence->confidence x 1/d, trunk–thigh angle	n	–2.46%	–6.58%	+4.12	0.005
LOO offset correction, thigh–shank angle	y*‡	+34.04%	–27.03%	+61.07	‡ withheld
LOO offset correction, frontal-plane thigh–shank	y*‡	+34.04%	–10.31%	+44.35	‡ withheld
confidence->confidence x 1/d, frontal-plane thigh–shank	n	–2.46%	–7.67%	+5.22	0.008
LOO offset correction, trunk–thigh angle	y*‡	+34.04%	–4.48%	+38.52	‡ withheld
criterion at matched $k = 2$, thigh–shank angle	n	+0.13%	–4.01%	+4.14	0.010
Gauss–Newton with rejection (w^2 -matched), thigh–shank angle	y	+3.45%	–7.45%	+10.90	0.010
no rejection → object-space rejection, thigh–shank angle	y	+0.51%	–6.09%	+6.60	0.012
GN: no rejection → reproj rejection, thigh–shank angle	y	+2.79%	–7.25%	+10.04	0.012
no rejection → object-space rejection, trunk–thigh angle	y	+0.51%	–5.83%	+6.34	0.014†
no rejection → reprojection rejection, thigh–shank angle	y	–0.10%	–8.16%	+8.07	0.024†
reprojection → object-space criterion, frontal-plane thigh–shank	y	+0.61%	–2.61%	+3.22	0.035†
uniform → confidence weighting, trunk–thigh angle	y	+3.10%	+5.69%	–2.59	0.071
DLT → Gauss–Newton (w^2 -matched), thigh–shank angle	y	–3.31%	+0.22%	–3.54	0.215
criterion at matched $k = 3$, thigh–shank angle	n	+0.58%	–1.74%	+2.31	0.238
uniform → confidence weighting, thigh–shank angle	y	+3.10%	+1.43%	+1.67	0.284
reprojection → object-space criterion, thigh–shank angle	y	+0.61%	+1.74%	–1.13	0.377
GN objective: w^2 -matched → unmatched (w), thigh–shank angle	y	–1.53%	–1.03%	–0.50	0.409
uniform → confidence weighting, frontal-plane thigh–shank	y	+3.10%	+3.06%	+0.04	0.947

A position-against-angle interaction cannot be run in native units, the two outcomes having different dimensions, so §2.5's both-scales rule is unenforceable for this family and every row here is relative-scale only. Where a row's two components share a sign its magnitude is denominator-dependent; where they oppose, the sign is not. All 20 interactions in the family are listed, survivors in bold. Thirteen of the twenty survive correction as the artefact prints it, but three of those are offset-table rows whose q is withheld (‡), so **ten interactions carry a valid verdict**, six of them on deployable arms; 1 that was significant unadjusted does not survive. Only ONE row pairs two individually surviving effects of opposite sign: an offset-table row cannot, because under the rule of Section 2.5 neither of its components has a valid q. That one row, Gauss–Newton rejection, itself survives the arbitrary-dependence sensitivity: it sits at rank 10 of 20 with $p = 0.0059$ against a Benjamini–Yekutieli threshold of 0.00695

there ($m = 20$, $H_m = 3.60$), and ten of the twenty survive. Its angle component carries a dagger. Section 3.2 states it that way. *dep.* marks an arm a validation study could deploy; **n** marks a diagnostic or control arm — the confidence $\times 1/d$ rows use ground-truth distances and the matched- k rows fix retention at a level the adaptive rule never reaches. ‡ marks a row on which no q is reported. Each participant's correction is built from the other ten, so the eleven differences are not exchangeable under the null and the sign-flip test does not apply; a rule declared before those rows were run required the q to be withheld. The only inferential statement available for them is a nested null on the [pos:HKA] gain, one-sided $p = 0.026$ (§2.5, Table S3e). *y** marks the population offset table, deployable in principle but computed here as a world-frame translation on trials whose headings span 56° across the study and $1.0\text{--}14.1^\circ$ within ten of the eleven participants, so this table does not transfer beyond this rig (§2.3). † marks a verdict that does not survive the arbitrary-dependence

sensitivity of Table S5. Exact 95% intervals for every surviving row are in Table S8. The 7 that do not are shown so that the family can be audited rather than summarised.

A surviving interaction is not by itself a dissociation.

Two rows pair component effects of opposite sign — LOO offset correction at the thigh–shank angle, and Gauss–Newton no rejection → reprojection rejection at the thigh–shank angle — but only the second has a valid q on BOTH components, because under the rule of Section 2.5 neither component of an offset-table row has one. One row therefore pairs two individually SURVIVING effects of opposite sign, and Section 3.2 reports that one. Elsewhere the two components share a sign (the confidence $\times 1/d$ rows, which harm both outcomes unequally) or one component is null on its own (the matched- k and no-rejection rows, whose position half does not move). Read each row's two component columns before reading its interaction.

Scale. Δ position and Δ angle are **means of per-participant ratios** — each participant's effect divided by that participant's own baseline error, then averaged — not ratios

Position and angle move apart, and a pooled position summary hides it




Fig. 1. A position gain and an angle change need not travel together, and the gap is what is tested. For each intervention the open marker is the 3D position effect on hip, knee and ankle and the filled marker is the joint-angle effect, both as a percentage of that arm's baseline error; the line is the gap between them and the right-hand number is the interaction (position minus angle, percentage points). Nine of the twenty declared interactions are drawn: every intervention that appears in a headline claim, at the angle that claim is made on. An asterisk marks $q < 0.05$ within the family of twenty; the offset-correction rows carry ‡ and no q , because Section 2.5 withholds it on leave-one-out rows. (a) Triangulation interventions. (b) The joint-centre offset correction, on its own scale.

The null appears to be a property of this rig as much as of the rule. On fixed subsets of the same ring both contrasts remained null at seven cameras but broke at five, where discarding bought 0.857 mm on the reprojection criterion and 0.428 mm on the object-space criterion; one subset per size is not an average, so this bounds the claim rather than measuring a five-camera rig (Table S23). Two controls also

of pooled means. The two differ only slightly once both use the baseline arm as denominator: the LOO position effect is +34.04% here against +9.848 mm on a 28.5 mm pooled baseline (34.6%).

The per-participant form is used because the interaction is a within-participant contrast and the inferential unit is the participant.

Under Gauss–Newton refinement, reprojection rejection improved lower-limb position by 0.858 mm while degrading the thigh–shank angle by 0.169° ($q = 0.006$ and $0.028^\dagger$), an interaction of 10.04 percentage points ($q = 0.012$; Fig. 1). Of the twenty declared interactions (Table S3i) this is the only one pairing two surviving effects of opposite sign, but its angle component, bounded within -0.016 to $+0.053^\circ$, does not survive the arbitrary-dependence sensitivity, so the dissociation is not demonstrated (Table S11). Gauss–Newton itself cost 1.384 mm relative to the direct linear transform.

cut against the criterion comparison. The object-space advantage fell from 4.62 mm to a non-significant 0.68 mm once lens distortion was corrected (Fig. 3b), and on the elbow chain, outside every declared family and under a subject-cluster bootstrap and not the exact test, reprojection rejection lowered angle error by 0.295° , the opposite sign to the other chains (Table S31).

D. Confidence weighting and the population offset correction

Confidence weighting improved lower-limb position by 3.10% relative to uniform weighting (Table 1), the frontal-plane thigh–shank angle by 0.078° ($q = 0.004$) and the trunk–thigh angle by 0.227° ($q = 0.006$). It did not detectably transmit to the thigh–shank angle itself (0.023° , $q = 0.647$) and neither interaction survived, so we claim no angle gain beyond position. The largest single effect in the study was at the peripheral joints, where confidence weighting bought

5.062 mm [+4.309, +5.930]. The offset alone, propagated through the angle definition, yielded 2.766° of mean thigh–shank error against 2.722° for the full pipeline on the same arm, so it appears sufficient on its own to produce an error of the observed size, although it is not the same error ($r = 0.199$; Tables S3e, S3f).

Table 1. 3D position error: the joint set changes the answer

(a) The same contrasts on two joint sets (mm)

contrast	dep.	ankles + wrists	hip, knee and ankle
uniform vs confidence weighting	y	+5.062 [+4.309, +5.930], $q = 0.004$	+0.917, $q = 0.004$
confidence vs confidence $\times 1/d$ weighting, no rejection	n	+0.256, $q = 0.038^\dagger$	−0.691, $q = 0.004$
reprojection vs object-space criterion, adaptive	y	+0.332, $q = 0.004$	+0.168, $q = 0.018^\dagger$
reprojection vs object-space criterion, matched $k = 1$	n	+0.148, $q = 0.009$	+0.029, n.s.
reprojection vs object-space criterion, matched $k = 2$	n	+0.428, $q = 0.004$	+0.039, n.s.
reprojection vs object-space criterion, matched $k = 3$	n	+0.782, $q = 0.004$	+0.169, n.s.
no rejection vs reprojection rejection, adaptive	y	+0.131, n.s.	−0.005, n.s.
no rejection vs object-space rejection, adaptive	y	+0.463, n.s.	+0.163, n.s.
uncorrected vs leave-one-participant-out offset correction	y*‡	—	+9.848, ‡ withheld
DLT vs Gauss–Newton, no rejection	y	—	−1.384, $q = 0.004$
Gauss–Newton: no rejection vs reprojection rejection	y	—	+0.858, $q = 0.006$

Bold marks a multiplicity-surviving effect in the hip/knee/ankle column, which is where the paper's claims about the joints an angle uses are made. *dep.* marks an arm a validation study could deploy; **n** marks a diagnostic or control arm that cannot carry a practical claim — the two distance weightings use ground-truth distances, and matched- k fixes retention at a level the adaptive rule never reaches. ‡ marks a row on which no q is reported. Each participant's correction is built from the other ten, so the eleven differences are not exchangeable under the null and the sign-flip test does not apply; a rule declared before those rows were run required the q to be withheld. The only inferential statement available for them is a nested null on the [pos:HKA] gain, one-sided $p = 0.026$ (§2.5, Table S3e). **y*** marks the population offset table, deployable in principle but computed here as a world-frame translation on trials whose headings span 56° across the study and 1.0 – 14.1° within ten of the eleven participants, so this table does not transfer

beyond this rig (§2.3). † marks a verdict that does not survive the arbitrary-dependence sensitivity of Table S5. Exact 95% intervals for every surviving row are in Table S8.

Each row's baseline is the arm it names first; for the no-rejection confidence arm that is 16.6 mm (ankles + wrists) and 28.5 mm (hip, knee, ankle), and the percentages in Table 1(c) use each row's own baseline rather than these two. The hips dominate the latter at 37.7–39.3 mm. Confidence $\times 1/d$ weighting reverses sign between the two joint sets, and the matched- k criterion advantage disappears entirely.

(b) **Decomposed by joint** (mm; exact sign-flip p , uncorrected, within this analysis; bold marks $p < 0.05$ among the per-joint entries, not a corrected verdict. The per-joint q -values are in Table S3b, where seven of the twelve survive and the bold ankle entry in the adaptive row does not; exact intervals are in Table S8)

contrast	hip	knee	ankle	three-joint mean
reprojection vs object-space, matched $k = 1$	+0.253 (0.084)	−0.322 (0.003)	+0.134 (0.148)	+0.022 (0.706)
reprojection vs object-space, matched $k = 2$	+0.530 (0.002)	−0.655 (0.002)	+0.219 (0.076)	+0.031 (0.654)
reprojection vs object-space, adaptive	−0.024 (0.835)	+0.339 (0.002)	+0.190 (0.035)	+0.168 (0.008)
uniform $\rightarrow$ confidence weighting	+1.564 (0.001)	+0.365 (0.001)	+0.822 (0.001)	+0.917 (0.001)

At matched retention the criterion improves the hip and degrades the knee, both significantly, and the three-joint mean cancels them to a null. Confidence weighting improves every joint. The mean column is the unweighted mean of the three per-joint effects and belongs to the within-limb per-subset family ($m = 20$), not to the $m = 50$ family of panel (a):

the two differ by ≤ 0.01 mm because panel (a) is computed over a 0.1% different observation set. Table S7 repeats every row here with the joint-centre offset removed.

(c) **Does the joint set change the effect?** — per participant, the effect on ankles and wrists minus the effect

on hip/knee/ankle, each as a percentage of that participant's of 8
own baseline at that joint set; BH-corrected within the family

contrast	Δ [pos:AW]	Δ [pos:HKA]	interaction	q
uniform vs confidence weighting	+23.41%	+3.10%	+20.31	0.002
criterion at matched $k = 2$	+2.54%	+0.13%	+2.41	0.002
criterion at matched $k = 3$	+4.49%	+0.58%	+3.91	0.002
confidence vs confidence $\times 1/d$ weighting	+1.46%	-2.46%	+3.91	0.002
criterion, adaptive rejection	+1.95%	+0.61%	+1.34	0.005
criterion at matched $k = 1$	+0.92%	+0.07%	+0.84	0.014
no rejection vs object-space rejection	+2.48%	+0.51%	+1.97	0.172
no rejection vs reprojection rejection	+0.54%	-0.10%	+0.64	0.618

All 8 are listed. 6 survive correction: all three matched-retention contrasts, which are controls; the confidence $\times 1/d$ contrast that carries the sign reversal; and two deployable ones, uniform versus confidence weighting ($q = 0.002$) and the adaptive criterion ($q = 0.005$). The 2 that do not are the rejection-versus-none contrasts, which are null on both joint sets. This interaction is run on the relative scale only, so by Section 2.5's rule it is scale-dependent; the within-limb control of panel (d), run on both scales, carries the claim.

(d) Within one limb: only the membership of the set changes. Eighty contrasts across three declared families, all corrected and all reported in Table S3n. They are not reproduced here: the thirty pairwise differences on each scale span ten free quantities rather than thirty, and after both declared sensitivities the panel reduces to a single deployable contrast — confidence weighting, {hip, knee, ankle} against {hip, ankle}, -1.05% and -0.276 mm ($q = 0.004$) — which Section 3.1 quotes in full. The families remain in the multiplicity accounting of Table S4b; only their presentation moves.

Correcting each participant with the mean offset of the other ten improved lower-limb position by 9.848 mm but worsened the thigh–shank angle by 0.606° $[-0.932, -0.287]$. No q is reported on either row, for the reason given in Section 2.5. A rule fixed in advance required the position gain to clear the whole of a nested null, which it did not, so whether the offset transfers across participants is not established here. Correcting with a participant's own offset, a ceiling because it needs participant-specific marker data, bought 1.094° at the trunk–thigh angle and nothing at the thigh–shank angle (0.166° , $p = 0.679$). Two features of the offset may explain why. Only 11.04 mm of the knee's 32.36 mm offset lies in the plane that moves flexion (Table 3), and the offset is not constant within a participant. Across quintiles of the reference thigh–shank angle its magnitude changed by +6.19 mm at the hip, -2.82 mm at the knee and -4.27 mm at the ankle (Fig. 4), so the largest posture dependence appears to be at the hip, on which the literature disagrees (Tang et al., 2022; Wade et al., 2023), although at $n = 11$ a correlation below about 0.6 would be undetectable, so neither account is refuted (Table S17).

IV. DISCUSSION

Within a participant, most of what the position metric measures at the hip and knee is a disagreement about where the joint centre lies and not an error of reconstruction, the offset accounting for 74–81% of squared error. Pooled across participants, as a cohort-level validation study measures it, the share is 47–57% at the hip and 61–70% at the knee. The

consequence for comparisons between pipeline variants was shown empirically, in that removing the offset moved per-joint verdicts. Das et al. (2023) and Pataky et al. (2013) have both shown a scalar summary over-simplifying a richer object; what this study adds is the size of the displacement relative to what a pipeline decision can change.

Whether a displacement of this size represents a difference of definition or an error of our own cannot be settled at the hip. The two candidate errors common to both arms, the synchronisation residual and the reference's own alignment error, are at most 15.8 mm at the hip and knee on the 95th-percentile bound and 3.62 mm respectively, against 27–37 mm. At the hip, however, the reference centre is itself a regression from external landmarks whose published accuracy, 26–33 mm at 95% certainty, is of the same size as the offset we report (Bell et al., 1989), whereas the knee and ankle references are marker midpoints (Tables S0, S3m). The hip offset can therefore not be apportioned between the two systems and no absolute accuracy claim follows. The hip is consequently both the joint at which putting the target back is least well defined and the joint carrying the largest reversal, which we cannot explain.

The offset was reproduced by three detectors in the same direction for every participant, and its mediolateral component at the hip is bilaterally symmetric and outward. Both observations are consistent with the interpretation of Needham et al. (2021), who attributed the hip and knee differences to the labelling conventions of the training data and not to the estimator. They are also consistent with the constant biases D'Haene et al. (2024) removed by subtracting a mean waveform difference before comparing systems.

This study has several limitations. Every markerless estimate was unfiltered against a 12 Hz-filtered reference, and on the one filtered subset a post-hoc test could not distinguish the angle cost from zero, which is the weaker reading we take (Table S18). The direct linear transform and Gauss–Newton comparison is confounded with a difference in objective, and the offset-removal analysis subtracts a mean offset that itself varies with posture, so what remains is not offset-free. The eleven participants are not a random eleven of the archive's fifteen, being three women short of the cohort it describes, and two of them shape what we report. One spans 53.7° of heading where the other ten span 1.0 – 14.1° and carries the largest thigh–shank error, so the bias-term argument of Table S14 assumes a constancy that participant does not have. Another sets the worst case in twelve of the fourteen headline contrasts, although no contrast depends on it. All results come from one rig, one

task and one laboratory, every headline contrast from a single detector, and transverse-plane kinematics were not assessed. The family structure was also exploratory, and boundaries chosen after seeing results are a degree of freedom no correction removes. Every angle effect we report lies below the smallest between-session minimal detectable change Wilken et al. (2012) give for sagittal knee angles, 4.16° to 6.39° across peaks, which may transfer as a scale but not as a threshold. Finally, a pipeline that solves inverse kinematics has no small defining point set, the analogous lever there being the marker set (Mundt et al., 2025), so we scope these results to direct triangulation. That excludes Uhlich et al. (2023) and Pagnon et al. (2021), which fit musculoskeletal models, along with the model-fitting systems dominating the reviews cited in Section 1.

V. CONCLUSION

On this nine-camera ring no per-joint triangulation effect at the hip, knee or ankle exceeded 2.5 mm and no single-step angle effect on the four chains exceeded 0.38° , as point estimates, with widest 95% limits of +3.417 mm and -0.583° . The offset correction is not a triangulation effect and moves 9.848 mm. The rejection claim rests on the narrower ± 0.35 mm at 90%, which was broken on a five-camera subset of the same ring. Re-expressed against each participant's own mean error, discarding worsened hip error by 1.726 mm (object-space) and 1.983 mm (reprojection), leaving the hip estimate noisier about its own mean and better centred on the reference at the same time. The ceiling is scoped to those four chains, since the two rejection contrasts move the shank-foot angle by -0.522° (reprojection) and -0.643° (object-space) at the toe, where the pre-declared veto returns no verdict because the heel reverses sign. Of the four chains, discarding cost two and reversed on the elbow, and where it cost, the direction is established and the size is not ($0.053-0.286^\circ$).

These results come from one rig and one task, most sit at the sign-flip floor where the test certifies direction and not magnitude, and the family structure is exploratory, so we offer them as hypotheses for confirmatory work. On that basis we suggest that per-joint position verdicts be read against the offsets they contain, that position be measured on the joints the reported angle uses, and that any claimed dissociation be tested as an interaction and not as a pair of separate verdicts.

Acknowledgements

We thank the authors of the BioCV dataset. Its collection was supported by CAMERA (EP/M023281/1) and CAMERA 2.0 (EP/T022523/1).

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During the preparation of this work the author(s) used Claude (Anthropic) in order to draft and revise manuscript text, to write and run the analysis and figure-generation code, and to check the manuscript's statements against the cited sources and the result files. After using this tool the author(s) reviewed and edited the content as needed and take(s) full responsibility for the content of the published article.

Data and code availability

The BioCV dataset is available from the University of Bath Research Data Archive (<https://doi.org/10.15125/BATH-01258>) on request under a data use agreement. All analysis code — triangulation arms, permutation and interval estimation, interaction testing, and figure generation — is available at **[repository DOI to be supplied on acceptance]**.

Conflict of interest statement

None declared.

Ethics

Secondary analysis of the BioCV dataset; no new data were collected and no additional approval was required. The original collection was approved by the Institutional Review Board of the University of Bath (EP1819052, 25/07/2019) and all participants gave written informed consent to participate and to data release.

[Funding and CRediT author contributions — to be completed.]

References

- Bell, A.L., Brand, R.A., Pedersen, D.R., 1989. Prediction of hip joint centre location from external landmarks. *Hum. Mov. Sci.* 8, 3–16.
- Benjamini, Y., Hochberg, Y., 1995. Controlling the false discovery rate: a practical and powerful approach to multiple testing. *J. R. Stat. Soc. B* 57, 289–300. doi:10.1111/j.2517-6161.1995.tb02031.x
- Benjamini, Y., Yekutieli, D., 2001. The control of the false discovery rate in multiple testing under dependency. *Ann. Stat.* 29, 1165–1188. doi:10.1214/aos/1013699998
- Brown, D.C., 1971. Close-range camera calibration. *Photogramm. Eng.* 37, 855–866.
- Cheng, X., Jiao, Y., Meiring, R.M., Sheng, B., Zhang, Y., 2025. Reliability and validity of current computer vision based motion capture systems in gait analysis: a systematic review. *Gait Posture* 120, 150–160. doi:10.1016/j.gaitpost.2025.04.016
- Das, K., de Paula Oliveira, T., Newell, J., 2023. Comparison of markerless and marker-based motion capture systems using 95% functional limits of agreement in a linear mixed-effects modelling framework. *Sci. Rep.* 13, 22880. doi:10.1038/s41598-023-49360-2
- D'Haene, M., Chorin, F., Colson, S.S., Guerin, O., Zory, R., Piche, E., 2024. Validation of a 3D markerless motion capture tool using multiple pose and depth estimations for quantitative gait analysis. *Sensors* 24, 7105. doi:10.3390/s24227105
- Della Croce, U., Leardini, A., Chiari, L., Cappozzo, A., 2005. Human movement analysis using stereophotogrammetry. Part 4: assessment of anatomical landmark misplacement and its effects on joint kinematics. *Gait Posture* 21, 226–237. doi:10.1016/j.gaitpost.2004.05.003
- Evans, M., Needham, L., Wade, L., Parsons, M., Colyer, S., McGuigan, P., Bilzon, J., Cosker, D., 2024. Synchronised video, motion capture and force plate dataset for validating markerless human movement analysis. *Sci. Data* 11, 1300. doi:10.1038/s41597-024-04077-3
- Gelman, A., Stern, H., 2006. The difference between "significant" and "not significant" is not itself statistically significant. *Am. Stat.* 60, 328–331. doi:10.1198/000313006X152649
- Ionescu, C., Papava, D., Olaru, V., Sminchisescu, C., 2014. Human3.6M: large scale datasets and predictive methods for 3D human sensing in natural environments. *IEEE Trans. Pattern Anal. Mach. Intell.* 36, 1325–1339. doi:10.1109/TPAMI.2013.248
- Jiang, T., Lu, P., Zhang, L., Ma, N., Han, R., Lyu, C., Li, Y., Chen, K., 2023. RTMPose: real-time multi-person pose estimation based on MMPose. *arXiv:2303.07399*. doi:10.48550/arXiv.2303.07399
- Kanko, R.M., Laende, E.K., Davis, E.M., Selbie, W.S., Deluzio, K.J., 2021. Concurrent assessment of gait kinematics using marker-based and markerless motion capture. *J. Biomech.* 127, 110665. doi:10.1016/j.jbiomech.2021.110665
- Karashchuk, P., Rupp, K.L., Dickinson, E.S., Walling-Bell, S., Sanders, E., Azim, E., Brunton, B.W., Tuthill, J.C., 2021. Anipose: a toolkit for robust markerless 3D pose estimation. *Cell Rep.* 36, 109730. doi:10.1016/j.celrep.2021.109730
- Lu, C.-P., Hager, G.D., Mjolsness, E., 2000. Fast and globally convergent pose estimation from video images. *IEEE Trans. Pattern Anal. Mach. Intell.* 22, 610–622. doi:10.1109/34.862199
- Mundt, M., Pagnon, D., Colyer, S., 2025. The influence of the marker set on inverse kinematics results to inform markerless motion capture annotations. *Sci. Rep.* 15, 14547. doi:10.1038/s41598-025-97219-5
- Needham, L., Evans, M., Cosker, D.P., Wade, L., McGuigan, P.M., Bilzon, J.L., Colyer, S.L., 2021. The accuracy of several pose estimation methods for 3D joint centre localisation. *Sci. Rep.* 11, 20673. doi:10.1038/s41598-021-00212-x
- Pagnon, D., Domalain, M., Reveret, L., 2021. Pose2Sim: an end-to-end workflow for 3D markerless sports kinematics—part 1: robustness. *Sensors* 21, 6530. doi:10.3390/s21196530
- Pagnon, D., Domalain, M., Reveret, L., 2022. Pose2Sim: an end-to-end workflow for 3D markerless sports kinematics—part 2: accuracy. *Sensors* 22, 2712. doi:10.3390/s22072712
- Pataky, T.C., Robinson, M.A., Vanrenterghem, J., 2013. Vector field statistical analysis of kinematic and force trajectories. *J. Biomech.* 46, 2394–2401. doi:10.1016/j.jbiomech.2013.07.031

- Tang, H., Pan, J., Munkasy, B., Duffy, K., Li, L., 2022. Comparison of lower extremity joint moment and power estimated by markerless and marker-based systems during treadmill running. *Bioengineering* 9, 574. doi:10.3390/bioengineering9100574
- Uhlrich, S.D., Falisse, A., Kidziński, Ł., Muccini, J., Ko, M., Chaudhari, A.S., Hicks, J.L., Delp, S.L., 2023. OpenCap: human movement dynamics from smartphone videos. *PLoS Comput. Biol.* 19, e1011462. doi:10.1371/journal.pcbi.1011462
- Varcin, F., Boocock, M.G., 2026. The accuracy, validity and reliability of Theia3D markerless motion capture for studying the biomechanics of human movement: a systematic review. *Artif. Intell. Med.* 173, 103332. doi:10.1016/j.artmed.2025.103332
- Wade, L., Needham, L., McGuigan, P., Bilzon, J., 2022. Applications and limitations of current markerless motion capture methods for clinical gait biomechanics. *PeerJ* 10, e12995. doi:10.7717/peerj.12995
- Wade, L., Needham, L., Evans, M., McGuigan, P., Colyer, S., Cosker, D., Bilzon, J., 2023. Examination of 2D frontal and sagittal markerless motion capture: implications for markerless applications. *PLoS ONE* 18, e0293917. doi:10.1371/journal.pone.0293917
- Wilken, J.M., Rodriguez, K.M., Brawner, M., Darter, B.J., 2012. Reliability and minimal detectable change values for gait kinematics and kinetics in healthy adults. *Gait Posture* 35, 301–307. doi:10.1016/j.gaitpost.2011.09.105

Figure captions